

CS6600 2024

Indian Institute of Technology, Madras

Authors

Sai Ashwin R NS24Z344

Date

October 20th 2024

IIT M 2024-25

Contents

1	Bimodal predictor	3
1.1	Question 1A	3
1.1.1	Trends in gcc_trace.txt	3
1.1.2	Trends in jpeg_trace.txt	3
1.1.3	Similarities	3
1.1.4	Differences	4
1.2	Question 1B	4
2	Gshare Predictor Results	5
2.1	Question 2A	5
2.1.1	Trends in gcc_trace.txt	5
2.1.2	Trends in jpeg_trace.txt	6
2.1.3	Similarities	6
2.1.4	Differences	6
2.2	Question 2B	6

1 Bimodal predictor

1.1 Question 1A

The following plots show the misprediction rate vs "m" for the 2 trace files:




Figure 1: Bimodal Predictor Misprediction Rate for different trace files

1.1.1 Trends in gcc_trace.txt

For the gcc benchmark, the misprediction rate decreases steadily as m increases from 7 to 12. The rate drops from 26.65% at $m = 7$ to 12.47% at $m = 12$. This trend indicates that increasing the number of low-order PC bits used to index the bimodal predictor significantly improves prediction accuracy for the gcc workload, particularly in the early range of m .

1.1.2 Trends in jpeg_trace.txt

For the jpeg benchmark, the misprediction rate shows only a slight improvement as m increases, starting at 7.92% for $m = 7$ and decreasing to 7.6% for $m = 12$. This smaller variation suggests that the jpeg trace is less sensitive to changes in the number of PC bits used for the bimodal predictor.

1.1.3 Similarities

Both benchmarks demonstrate a general decrease in misprediction rate as m increases, showing that using more PC bits for indexing improves the accuracy of the bimodal predictor. Additionally, the improvement in misprediction rate flattens out for higher values of m , suggesting diminishing returns as m approaches 12.

1.1.4 Differences

The gcc benchmark experiences a much larger reduction in misprediction rate (over 14% decrease from $m = 7$ to $m = 12$), while the jpeg benchmark shows only a minor reduction (around 0.3%). This indicates that the bimodal predictor performs much better with increasing m for the gcc workload, while the jpeg workload remains relatively unaffected by changes in m .

In addition, the gcc benchmark on average experiences a **much larger misprediction rate** than the jpeg benchmark with the maximum misprediction rate being 7.92 % for the jpeg trace vs the 26.65 % of the gcc trace. This suggests that the gcc benchmark has more **complicated control flow logic** than the jpeg benchmark. This can be reasoned easily by noting that gcc is a compiler, while the jpeg trace is probably that of a simple jpeg parser. Thus, the gcc benchmark might need a more sophisticated branch predictor.

1.2 Question 1B

We are given a maximum budget of 16 kB for the branch predictor. The amount of storage required when there are " m " bits of the PC used is $2^m * 2 = 2^{m+1}$ bits. Thus,

$$\begin{aligned}2^{m+1} &= 16 * 1024 * 8 \\2^{m+1} &= 2^{17} \\m &= 16\end{aligned}$$

For the **gcc_trace.txt** benchmark, the misprediction rate decreases as m increases from 7 to 12:

- At $m = 7$, the misprediction rate is 26.65%.
- At $m = 10$, it drops to 15.67%.
- At $m = 11$, it drops to 13.65%
- At $m = 12$, the misprediction rate further decreases to 12.47%.

It is clear that the reduction slows down as m increases beyond 10.

Therefore, for **gcc_trace.txt**, we can choose $m = 11$ as the optimal value, as the reduction in misprediction rate beyond this point is close to only 1%, whereas the storage required grows exponentially.

For the **jpeg_trace.txt** benchmark, the misprediction rate also decreases as m increases, but the reduction is less significant:

- At $m = 7$, the misprediction rate is 7.92%.
- At $m = 9$, it reduces to 7.74%.

- At $m = 10$, it reduces to 7.7%.
- By $m = 11$, it only decreases slightly to 7.6%.

Since the misprediction rate improvement is very small after $m = 9$, the best choice is $m = 9$.

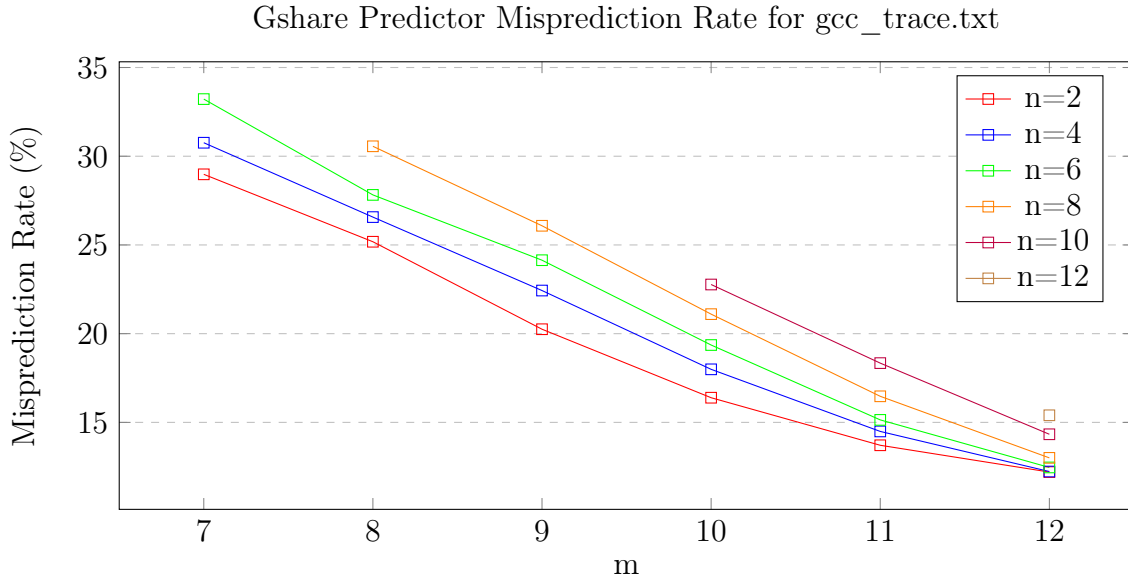
For both configurations, the optimal value chosen by us is less than the limit of $m = 16$.

In conclusion, the optimal design for **gcc_trace.txt** is **$m = 11$** , while the optimal design for **jpeg_trace.txt** is **$m = 9$** . Both designs minimize misprediction rates while staying within the 16 kB storage limit.

2 Gshare Predictor Results

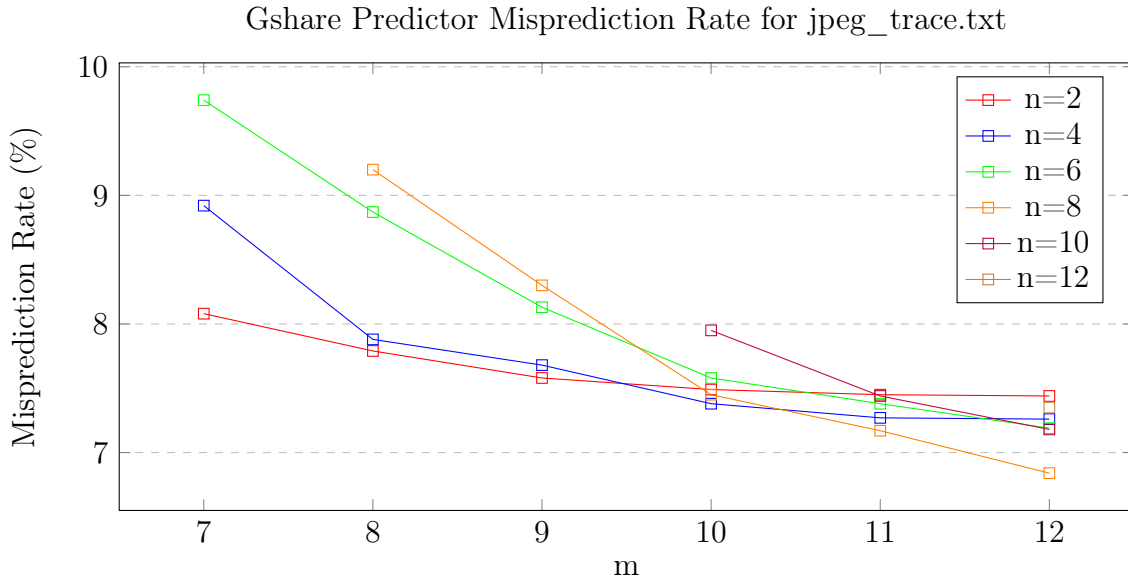
2.1 Question 2A

The following plots show the misprediction rate vs " m " for different values of " n " for the 2 trace files.



2.1.1 Trends in gcc_trace.txt

For the gcc benchmark, the misprediction rate decreases as the value of m increases from 7 to 12, across all values of n . This shows that increasing the number of PC bits used for indexing (m) improves the prediction accuracy. The impact of n (global history register size) is more pronounced for smaller m values. However, as m increases, the differences in misprediction rates between different values of n diminish, converging around $m = 12$.



2.1.2 Trends in jpeg_trace.txt

For the jpeg benchmark, the misprediction rate also decreases as m increases, but the decrease is more gradual compared to the gcc benchmark. The impact of n on the misprediction rate is less significant in this case, with a very small difference in rates across different values of n . This indicates that the jpeg workload is less sensitive to the global history register size and relies more on increasing m .

2.1.3 Similarities

Both benchmarks show a consistent decrease in misprediction rate as m increases. In addition, the misprediction rates tend to converge as m approaches 12, implying that at higher values of m , the impact of n becomes less important in both benchmarks.

2.1.4 Differences

The gcc benchmark exhibits higher misprediction rates overall (ranging from 12% to 35%) compared to the jpeg benchmark (6.8% to 10%). Furthermore, gcc is more sensitive to changes in n , particularly for smaller values of m . In contrast, the jpeg benchmark shows minimal sensitivity to n , indicating that increasing m is more impactful for improving prediction accuracy in the jpeg benchmark.

2.2 Question 2B

We are given a storage budget of 16 kB. We can write the number of bits required in terms of m and n :

$$\begin{aligned} \text{Storage} &= n + 2^m * 2 \\ &= n + 2^{m+1} \end{aligned}$$

For the **gcc_trace.txt** benchmark, the misprediction rate decreases significantly as m increases. For instance, for $n = 2$, the misprediction rate decreases from 28.98% at $m = 7$ to 12.2% at $m = 12$. However, as m increases beyond 10, the reduction in misprediction rate becomes less pronounced. Additionally, increasing n actually increases the misprediction rate, rather counter intuitively. This could be attributed to the **inertia** introduced by increasing the number of bits in the BHR. Therefore, for **gcc_trace.txt**, the optimal design is $\mathbf{m} = 11$ and $\mathbf{n} = 2$.

For the **jpeg_trace.txt** benchmark, the misprediction rate is less sensitive to increases in m and n . Most of the configurations seem to saturate at just above 7% after $m = 10$. We see diminishing returns after $m = 10$, and the best performing configuration at $m = 10$ has $n = 4$. Increasing n has a negligible effect on storage, while storage increases exponentially with m . Therefore, for **jpeg_trace.txt**, the optimal design is $\mathbf{m} = 10$ and $\mathbf{n} = 4$.

In conclusion, the optimal design for **gcc_trace.txt** is $\mathbf{m} = 11$ and $\mathbf{n} = 2$, while the optimal design for **jpeg_trace.txt** is $\mathbf{m} = 10$ and $\mathbf{n} = 4$. Both configurations minimize misprediction rates while respecting the storage constraint of 16 kB.